

A Blockchain Consensus Protocol With Horizontal Scalability

Kelong Cong
kelong.cong@epfl.ch
EPFL

Zhijie Ren
z.ren@tudelft.nl
TU Delft

Johan Pouwelse
peer2peer@gmail.com
TU Delft

IFIP Networking, 2018

Outline

Introduction

- The dangers of centralisation

- Related work

- Research question

System architecture

- System model

- Architecture overview

- Extended TrustChain

- Consensus protocol

- Transaction protocol

- Validation protocol

Experimental results

Conclusion

Outline

Introduction

- The dangers of centralisation

- Related work

- Research question

System architecture

- System model

- Architecture overview

- Extended TrustChain

- Consensus protocol

- Transaction protocol


- Validation protocol

Experimental results

Conclusion

The dangers of centralisation

- ▶ Technological advancements give us convenience
- ▶ But it puts central authorities in control
- ▶ Most are motivated exclusively by profit
- ▶ Not always in the interest of the “users”¹

¹Typically users of some free service X are, in fact, used by X. 

Blockchain: a new hope?

- ▶ Blockchains are distributed (replicated) ledgers with no central control
- ▶ They enable internet-scale consensus for the first time
- ▶ Some initial applications include:
 - ▶ Digital cash (e.g., Bitcoin, Litecoin)
 - ▶ Domain name system (e.g., Namecoin)
 - ▶ Storage rental (e.g., Filecoin)
 - ▶ General purpose (e.g., Ethereum)

A Blockchain Consensus Protocol With Horizontal Scalability

└ Introduction

└ The dangers of centralisation

└ Blockchain: a new hope?

- Blockchains are distributed (replicated) ledgers with no central control
- They enable internet-scale consensus for the first time
- Some initial applications include:
 - Digital cash (e.g., Bitcoin, Litecoin)
 - Domain name system (e.g., Namecoin)
 - Storage rental (e.g., Filecoin)
 - General purpose (e.g., Ethereum)

- Explain blockchain systems—consensus

Blockchain: not there yet

- ▶ All blockchain systems have a consensus algorithm
- ▶ Early consensus algorithms (PoW) do not scale
- ▶ Bitcoin is limited to 7 transactions per second
- ▶ 100,000 transaction backlog in May 2017
- ▶ We require horizontal scalability for ubiquitous use
- ▶ More users → more transactions per second globally

Related work

- ▶ Off-chain solution
 - ▶ Lightning Network
 - ▶ Perun
- ▶ On-chain solution
 - ▶ Parameter tuning
 - ▶ BFT consensus (e.g. Tendermint, ByzCoin)
 - ▶ Sharding (e.g. Elastico, OmniLedger)

Related work

State-of-the-art—Sharding:

- ▶ Split state into multiple shards
- ▶ Shards run consensus algorithm in parallel

Challenges:

- ▶ Choosing and evolving the shard size
- ▶ Perform atomic inter-shard transactions
- ▶ Parameter choice highly depends on the application

Research question

How can we design a *blockchain consensus protocol* that is *fault-tolerant*, *horizontally scalable*, and able to reach *global consensus*?

- ▶ Blockchain consensus protocol—application neutral, e.g., PoW
- ▶ Fault-tolerant—tolerate a number of malicious nodes
- ▶ Horizontal scalability—more nodes in the network leads to higher transaction throughput
- ▶ Global consensus—all node should agree on a global state

Outline

Introduction

- The dangers of centralisation

- Related work

- Research question

System architecture

- System model

- Architecture overview

- Extended TrustChain

- Consensus protocol

- Transaction protocol


- Validation protocol

Experimental results


Conclusion

Intuition and idea explored in this thesis

- ▶ It is expensive to verify and reach consensus on all transactions
- ▶ Our idea: we decouple consensus and validation
- ▶ A single digest represents an arbitrarily large number of transactions
- ▶ Reach consensus on the small digest
- ▶ Nodes then independently check the validity of the transactions of interest



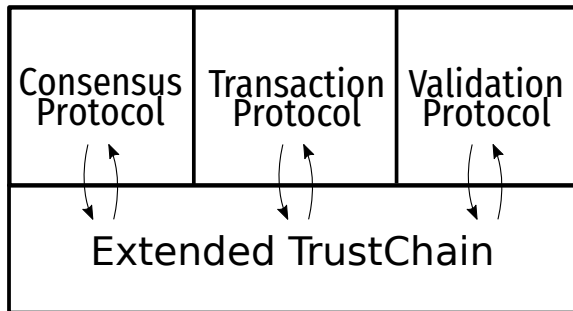
Early blockchains



Our idea

Architecture overview

The four components of CHECO




A Blockchain Consensus Protocol With Horizontal Scalability

System architecture

Architecture overview


Architecture overview

The four components of CHECO



- The primary data structure is the Extended TrustChain, extension of our prior work
- The three protocols the tasks as their name suggests
- They are independent and run concurrently
- The only synchronisation happens via the Extended TrustChain
- But in no part of those protocol do we lock the Extended TrustChain

Extended TrustChain



2018-04-23


A Blockchain Consensus Protocol With Horizontal Scalability

System architecture

Extended TrustChain

Extended TrustChain


Extended TrustChain



- In this example there are three nodes
- Each node maintains their personal hash chain and genesis block
- Squares are TX blocks and circles are CP blocks
- Explain the block content in caption
- The dotted line represent pairs of TX blocks
- Geared with the understanding of our data structure, we are ready to talk about the consensus protocol


Extended TrustChain: Transaction (TX) block

- ▶ Goal: record transactions
- ▶ A transaction is represented by a pair of TX blocks, i.e. a contract signed by both parties




Extended TrustChain: Checkpoint (CP) block

- ▶ Goal: represent the state of the chain using a single digest
- ▶ A collection of CP blocks from all the nodes represent the state of the system
- ▶ Nodes become aware of the system state by running our consensus protocol




Extended TrustChain: Fragment of a TX block




Consensus protocol

- ▶ Goal 1: reach consensus on a collection of CP blocks amongst all the nodes
- ▶ Goal 2: create new CP blocks at the end of the protocol
- ▶ Uses an existing fault-tolerant consensus algorithm (HoneyBadgerBFT [1]) as the building block
- ▶ But it cannot be used in a large network due to high communication complexity
- ▶ We overcome this limitation by selecting a small number of *facilitators* from the network to run HoneyBadgerBFT


Consensus protocol






collect CP blocks



disseminate result



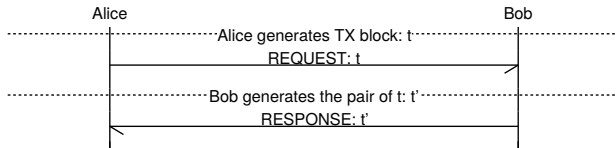
create new CP block
elect new facilitators



Consensus protocol: properties

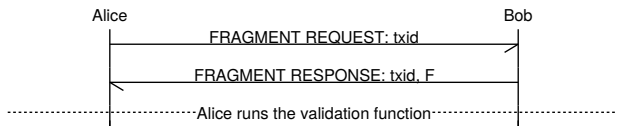
- ▶ *Agreement*: Every correct outputs the same set of facilitators.
- ▶ *Validity*: The consensus result is valid such that a new set of facilitators can be computed from it.
- ▶ *Termination*: Every correct node eventually outputs a set of facilitators.

Transaction protocol



- ▶ Two TX blocks are generated on the chains of Alice and Bob
- ▶ No guarantee that nodes follow this protocol

Validation protocol



- ▶ To check that the transaction protocol is correctly followed
- ▶ Alice needs the fragment of the TX on Bob's hash chain
- ▶ Validation function checks whether the fragment is OK and contain the transaction
- ▶ Can be generalised—any node may run the validation protocol on any transaction (does not need to be their own)

Validation protocol: properties

Consensus on CP blocks → consensus on transactions

- ▶ CP blocks of the fragments are “anchored” due to the consensus protocol
- ▶ It is difficult to modify the fragment once “anchored”
- ▶ Since the transaction protocol and the validation protocol only use point-to-point communication, we achieve horizontal scalability.

Outline

Introduction

- The dangers of centralisation

- Related work

- Research question

System architecture

- System model

- Architecture overview

- Extended TrustChain

- Consensus protocol

- Transaction protocol

- Validation protocol

Experimental results


Conclusion

Implementation and experiment setup


- ▶ Free and open source implementation on Github:
<https://github.com/kc1212/checo>
- ▶ SHA256 for hash functions and Ed25519 for digital signature
- ▶ Experiment on the DAS-5²
- ▶ Up to 1500 nodes

²<http://www.cs.vu.nl/das5/>

Validated transaction throughput (random node)



Validated transaction throughput (fixed neighbour)




2018-04-23

A Blockchain Consensus Protocol With Horizontal Scalability

└ Experimental results


└ Validated transaction throughput (fixed

Validated transaction throughput (fixed neighbour)



No need to request for fragment every time a TX needs to be validated.
Upon receiving a fragment, validate as many TX as possible.

Stress test (fixed neighbour)



Outline

Introduction

The dangers of centralisation

Related work

Research question

System architecture

System model

Architecture overview

Extended TrustChain

Consensus protocol

Transaction protocol

Validation protocol

Experimental results

Conclusion

Conclusion

Our work answers the research question.

How can we design a *blockchain consensus protocol* that is *fault-tolerant*, *horizontally-scalable*, and able to reach *global consensus*?

- ▶ Fault-tolerance is achieved using HoneyBadgerBFT
- ▶ Horizontal-scalability is achieved by separating consensus and validation, demonstrated experimentally
- ▶ Global-consensus on transactions is achieved via consensus on CP blocks

Bibliography

- [1] A. Miller, Y. Xia, K. Croman, E. Shi, and D. Song, “The honey badger of bft protocols”, in *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security*, ACM, 2016, pp. 31–42.

Thank you

Any questions?

TX block

1. Hash pointer to the previous block
2. Sequence number
3. Transaction ID
4. Public key of the counterparty
5. Transaction message m
6. Signature the five items above

A transaction is represented by a *pair* of TX blocks


CP block

1. Hash pointer to the previous block
2. Sequence number
3. Digest of consensus result, i.e. a set of CP blocks
4. Round number r
5. Signature on the four items above

Background on ACS

- ▶ Asynchronous common subset
- ▶ A simplification of HoneyBadgerBFT [1]
- ▶ n nodes
- ▶ t nodes may be malicious
- ▶ Input: every node proposes a set of values, e.g., $\{A, B\}, \{B, C\}, \dots$
- ▶ Output: set union of the majority, e.g., $\{A, B, C, \dots\}$

Consensus protocol: part 1




2018-04-23

A Blockchain Consensus Protocol With Horizontal Scalability


- Extras
 - Consensus protocol
 - Consensus protocol: part 1

Consensus protocol: part 1



Suppose we are in a state where C_{r-1} has just been agreed by some facilitators but not yet propagated.

Consensus protocol: part 2



2018-04-23


A Blockchain Consensus Protocol With Horizontal Scalability

└ Extras

└ Consensus protocol

└ Consensus protocol: part 2

Consensus protocol: part 2





Nodes receive consensus result \mathcal{C}_{r-1} , first n nodes ordered by $H(\mathcal{C}_{r-1}||pk)$ become \mathcal{F}_{r-1} , send the new CP blocks to \mathcal{F}_{r-1} .

Consensus protocol: part 3


facilitators: lottery
 $\{c_{a,3}, c_{b,0}, c_{c,4}, c_{d,3}\}$

round: 4

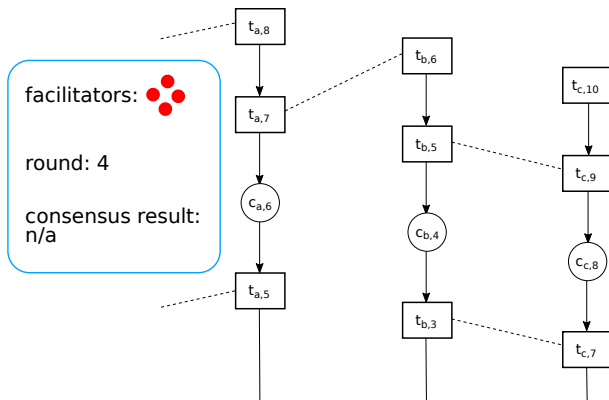
consensus result:
n/a



Consensus protocol: part 4



Consensus protocol: part 5



2018-04-23


A Blockchain Consensus Protocol With Horizontal Scalability

Extras

Consensus protocol


Consensus protocol: part 5

Consensus protocol: part 5



Transactions carry on as usual in round r , while facilitators are trying to reach consensus on the new CP blocks concurrently.

Consensus protocol: part 6



2018-04-23


A Blockchain Consensus Protocol With Horizontal Scalability

└ Extras

└ Consensus protocol


└ Consensus protocol: part 6

Consensus protocol: part 6




\mathcal{F}_{r-1} agree and disseminate \mathcal{C}_r , CP blocks at round $r - 1$ ($c_{a,6}, c_{b,4}, c_{c,8}$) should be in \mathcal{C}_r .

Effect of the number of facilitators (fixed neighbours)



Effect of the number of facilitators (random neighbours)



Future work

- ▶ Implement and experiment with a concrete application
- ▶ Analyse the system in the permissionless environment
- ▶ Improve fault tolerance