

Blockchain Consensus Protocol with Horizontal Scalability

K. Cong

Faculty of Electrical Engineering, Mathematics and Computer Science
Delft University of Technology

July 26, 2017

Outline

Introduction

- Motivation

- Research question

- Intuition

System architecture

- System model

- Architecture overview

- Extended TrustChain

- Consensus protocol

- Transaction protocol

- Validation protocol

Analysis of correctness and performance

- Correctness of the consensus protocol

- Correctness of the validation protocol

- Linear global throughput argument

Experimental results

Conclusion

Outline

Introduction

- Motivation

- Research question

- Intuition

System architecture

- System model

- Architecture overview

- Extended TrustChain

- Consensus protocol

- Transaction protocol

- Validation protocol

Analysis of correctness and performance

- Correctness of the consensus protocol

- Correctness of the validation protocol

- Linear global throughput argument

Experimental results

Conclusion

Motivation

- ▶ Blockchain systems offer an alternative to central authorities for the first time
- ▶ Market cap and trade volume figures indicate they are here to stay
- ▶ Early blockchain systems are not scalable (7 TX/s for Bitcoin)
- ▶ Parameter tuning leads to centralisation

Research question

How do we design a *blockchain consensus protocol* that is *fault tolerant*, *scalable* and can reach *global consensus*?

Intuition

- ▶ A restaurant owner does not report all of its transactions with an central authority
- ▶ Ocassionally a customer may leave without paying and this event is reported to a central authority
- ▶ Our blockchain system achieves scalability using the same idea

Outline

Introduction

Motivation

Research question

Intuition

System architecture

System model

Architecture overview

Extended TrustChain

Consensus protocol

Transaction protocol

Validation protocol

Analysis of correctness and performance

Correctness of the consensus protocol

Correctness of the validation protocol

Linear global throughput argument

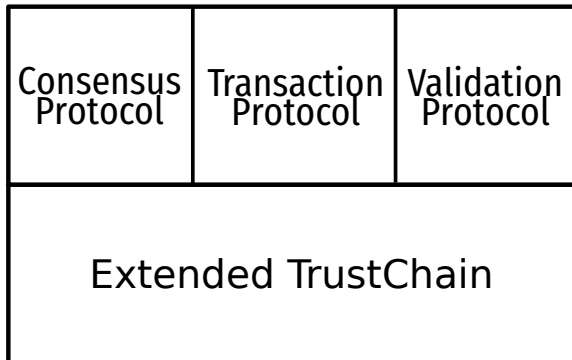
Experimental results

Conclusion

System model

- ▶ Population size is N
- ▶ n nodes are facilitators, t nodes are malicious (Byzantine)
- ▶ $n \geq 3t + 1$
- ▶ $N \geq n + t$
- ▶ Purely asynchronous channels with eventual delivery
- ▶ Public key infrastructure
- ▶ Random oracle model
- ▶ Application neutral

Architecture overview



Extended TrustChain

- ▶ Everyone has their own chain
- ▶ Two types of blocks, transaction (TX) blocks and checkpoint (CP) blocks
- ▶ A transaction involves two parties and results in two TX blocks
- ▶ A CP block captures the chain state
- ▶ TX and CP blocks are chained together using hash pointers

Extended TrustChain




Figure: TX block is a six-tuple: $t_{u,i} = \langle H(b_{u,i-1}), i, txid, pk_v, m, sig_u \rangle$.
 CP block is a five-tuple: $c_{u,i} = \langle H(b_{u,i-1}), i, H(C_r), r, sig_u \rangle$.

Consensus protocol overview

1. In round r , n out of N lucky nodes are selected at random to act as facilitators
2. Facilitators run a BFT (Byzantine Fault Tolerant) consensus algorithm to agree on a set of CP blocks
3. Disseminate the consensus result \mathcal{C}_r , i.e. the CP blocks
4. Repeat

Consensus protocol properties

$\forall r \in \mathbb{N}$, the following properties must hold.

- ▶ *Agreement*: If one correct node outputs a list of facilitators \mathcal{F}_r , then every node outputs \mathcal{F}_r
- ▶ *Validity*: If any correct node outputs \mathcal{F}_r , then
 1. $|\mathcal{C}_r| \geq N - t$ must hold for the \mathcal{C}_r which was used to create \mathcal{F}_r ,
 2. \mathcal{F}_r must contain at least $n - t$ honest nodes and
 3. $|\mathcal{F}_r| = n$.
- ▶ *Fairness*: Every node with a CP block in \mathcal{C}_r should have an equal probability of becoming a member of \mathcal{F}_r .
- ▶ *Termination*: Every correct node eventually outputs some \mathcal{F}_r .

Consensus protocol




Figure: Suppose we are in a state where \mathcal{C}_{r-1} has just been agreed by some facilitators but not yet propagated.

Consensus protocol




Figure: Nodes receive consensus result C_{r-1} , compute the new facilitators \mathcal{F}_{r-1} , and send out the new CP blocks.

Consensus protocol




Figure: Transactions carry on as usual in round r , while facilitators are trying to reach consensus on the new CP blocks concurrently.

Consensus protocol

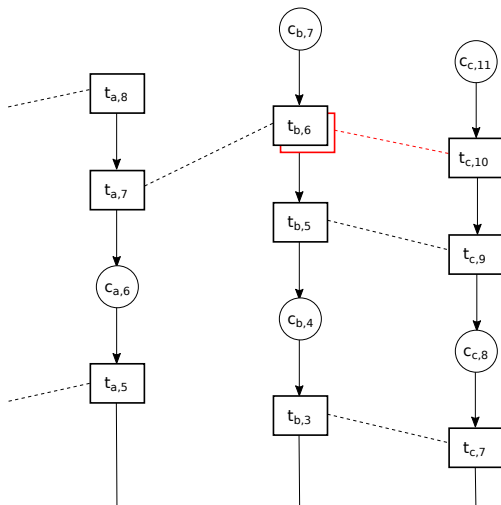


Figure: CP blocks at round $r - 1$ should be in \mathcal{C}_r . The first n nodes ordered by $H(\mathcal{C}_r || pk)$ become \mathcal{F}_r .

Transaction protocol

- ▶ Request (`tx_req`) and response (`tx_resp`) protocol
- ▶ Two TX blocks containing the same *txid* are generated
- ▶ Non-blocking

Transaction protocol

Create $t_{u,h}$ and send $\langle \text{tx_req}, t_{u,h} \rangle$ to start a transaction.

Upon $\langle \text{tx_req}, t_{v,j} \rangle$ from v

$\langle -, -, txid, pk_v, m, - \rangle \leftarrow t_{v,j}$

▷ unpack $t_{v,j}$

$\text{new_tx}(pk_u, m, txid)$

▷ create and store $t_{u,i}$

store $t_{v,j}$ as the pair of $t_{u,h}$

send $\langle \text{tx_resp}, t_{u,h} \rangle$ to v

Upon $\langle \text{tx_resp}, t_{v,j} \rangle$ from v

$\langle -, -, txid, pk_v, m, - \rangle \leftarrow t_{v,j}$

▷ unpack $t_{v,j}$

store $t_{v,j}$ as the pair of the TX with identifier $txid$

Validation protocol overview

- ▶ Request (`vd_req`) and response (`vd_resp`) protocol
- ▶ Transactions are in three states—*valid*, *invalid* and *unknown*
- ▶ The goal is to identify which state a given transaction is in

Validation protocol properties

- ▶ *Correctness*: The validation protocol outputs the correct result according to the aforementioned validity definition.
- ▶ *Agreement*: If any correct node decides on the validity (except when it is *unknown*) of a transaction, then all other correct nodes are able to reach the same conclusion or *unknown*.
- ▶ *Liveness*: Any valid transactions can be validated eventually.

Validity definition

Function `get_validity()` validates the transaction $t_{u,i}$

Check that v sent the correct agreed fragment, otherwise return *unknown*.

$\langle -, -, txid, pk_v, m, - \rangle \leftarrow t_{u,i}$

if number of blocks of $txid$ in $F_{v,j} \neq 1$ **then**

return *invalid*

▷ TX exists

$\langle -, -, txid', pk'_u, m', - \rangle \leftarrow t_{v,j}$

if $m \neq m' \vee pk_u \neq pk'_u$ **then**

return *invalid*

▷ no tampering

return *valid*

Validation protocol

Send $\langle \text{vd_req}, txid \rangle$ to v to begin.

Upon $\langle \text{vd_req}, txid \rangle$ from v

$t_{u,i} \leftarrow$ the transaction identified by $txid$

$F_{u,i} \leftarrow \text{agreed_fragment}(t_{u,i})$

send $\langle \text{vd_resp}, txid, F_{u,i} \rangle$ to v

Upon $\langle \text{vd_resp}, txid, F_{v,j} \rangle$ from v

$t_{u,i} \leftarrow$ the transaction identified by $txid$

set the validity of $t_{u,i}$ to $\text{get_validity}(t_{u,i}, F_{v,j})$

Outline

Introduction

Motivation

Research question

Intuition

System architecture

System model

Architecture overview

Extended TrustChain

Consensus protocol

Transaction protocol

Validation protocol

Analysis of correctness and performance

Correctness of the consensus protocol

Correctness of the validation protocol

Linear global throughput argument

Experimental results

Conclusion

Correctness of the consensus protocol

Theorem

For all rounds, the consensus protocol satisfies agreement, validity, fairness and termination.

Proof.

(sketch) Because consensus result are eventually delivered and the properties of ACS, we get agreement, validity and termination. Fairness is from the fact that we model $H(\cdot)$ as a random oracle (RO) and the input to the RO is different for every node, thus the list of nodes ordered by $H(C_r || pk)$ is a random permutation of those nodes. □

Correctness of the validation protocol

Theorem

If any correct node decides on the validity (except when it is unknown) of a transaction, then all other correct nodes are able to reach the same conclusion or unknown.

Proof.

(sketch) Proof by contradiction. For this attack to work, the adversary must be able to create two different fragments but with the same checkpoint enclosure. We model $H(\cdot)$ as a RO, so the adversary need to query the RO a exponential¹ number of times. But the adversary can only query the RO a polynomial number of times. □

¹In terms of the security parameter.

Linear global throughput argument

- ▶ Throughput has a notion of time but our model is asynchronous
- ▶ Additional assumptions needed to make the argument—every unit of communication takes a non-negligible amount of time to process
- ▶ Bandwidth relation— $NC \geq r_{tx}I$, where I is $O(N)$
- ▶ If r_{tx} satisfies the inequality, then LHS and RHS grows at the same rate, thus we have linear global throughput

Outline

Introduction

Motivation

Research question

Intuition

System architecture

System model

Architecture overview

Extended TrustChain

Consensus protocol

Transaction protocol

Validation protocol

Analysis of correctness and performance

Correctness of the consensus protocol

Correctness of the validation protocol

Linear global throughput argument

Experimental results

Conclusion


Implementation and experiment setup

- ▶ Prototype implementation on Github²
- ▶ SHA256 for hash functions and Ed25519 for digital signature
- ▶ Experiment on the DAS-5³


²<https://github.com/kc1212/consensus-thesis-code>

³<http://www.cs.vu.nl/das5/>

Throughput vs population size (random neighbour)



Throughput vs population size (fixed neighbour)



Outline

Introduction

- Motivation

- Research question

- Intuition

System architecture

- System model

- Architecture overview

- Extended TrustChain

- Consensus protocol

- Transaction protocol

- Validation protocol

Analysis of correctness and performance

- Correctness of the consensus protocol

- Correctness of the validation protocol

- Linear global throughput argument

Experimental results

Conclusion

Conclusion

How do we design a *blockchain consensus protocol* that is *fault tolerant*, *scalable* and can reach *global consensus*?

- ✓ Fault tolerant up to t nodes
- ✓ Horizontal scalability
- ✓ Global consensus

Future work

- ▶ Improve fault tolerance
- ▶ Improve fork detection
- ▶ Analyse the system in the permissionless environment
- ▶ Concrete application