TODO TITLE

TODO AUTHOR

TODO TITLE

Master's Thesis in Computer Science

Parallel and Distributed Systems group
Faculty of Electrical Engineering, Mathematics, and Computer Science
Delft University of Technology

TODO AUTHOR

16th June 2017

Abstract

TODO ABSTRACT

Preface

TODO MOTIVATION FOR RESEARCH TOPIC

TODO ACKNOWLEDGEMENTS

TODO AUTHOR

Delft, The Netherlands
16th June 2017

Contents

Chapter 1

Introduction

Blockchains is one of the most prominent but also controversial technologies of our time. Never before did we have an incorruptable ledger that can be shared with mutually untrustful parties. The ledger in fact can contain any type of data. Thus enabling applications in myriad of fields, such as finance [**todo**], heath care [**todo**], logistics [**todo**], energy [**todo**] and so on. However, not everyone is convinced on the practicality of blockchain systems. Bram Cohen, the inventor of BitTorrent said "It's getting a lot more media attention than the actual impact it's having so far." [5] Forbes listed a few reasons on why we should be skeptical about blockchains [2], some of them include its accessibility for the layperson and its scalability issues.

Scalability is indeed one of the key challenges we face in blockchain systems today. Bitcoin [**bitcoin**], the largest permissionless[1] blockchain system in terms of market capitalisation [**bitcoinmarketcap**] has a maximum transaction rate of merely 7 transaction per second (TX/s). This is due to the consensus mechanism in Bitcoin, namely proof-of-work (PoW), miners can only create new blocks every 10 minutes and every block cannot be larger than 1 megabyte. Payment processors in use today such as Visa can handle transaction rates in the order of thousands [8]. While Bitcoin may be a revolutionary phenomenon, it clearly cannot be ubiquitous in its current state.

An different approach is to not reach global consensus at all. For instance in TrustChain [**multichain**] and Tangle [**tangle**], nodes in the network only store their personal ledger. Since consensus is left out, nodes can perform transactions as fast as their machine and network allows. The downside of this approach is that it cannot prevent fraud (it is possible to detect fraud). To examplify, a malicious node Mallory may claim she has 3 units of currency to Alice, but in reality Mallory already spent all of it on Bob. If there is no global consensus and Bob and Alice never communicate, then

---

[1]Explain permissionless

the 3 units that Alice is about to receive is nonexistent.

The scalability property of TrustChain and Tangle are exceptionally desirable. The global consensus mechanism of Bitcoin and many other blockchain systems are also worthwhile for detecting or preventing fraud. These two properties may seem mutually exclusive, but in this work, we demonstrate the opposite. Specifically, we answer the following research question in the affirmative. *Is it possible to design a blockchain fabric that can reach global consensus on the state of the system and also scalable?* We define scalability as a property where if more nodes join the system, then the transaction rate should increases.

Our primary insight came from observing the differences between how transactional systems work in the real world and how they work in blockchain systems like Bitcoin. Take a restaurant owner for example, most of the time the customer is honest and pays the bill. There is no need for the customer or the restaurant owner to report the transaction to any central authority because both parties are happy with the transaction. On the other hand, if the customer leaves without paying the bill, then the restaurant owner would report the incident to some central authority, e.g. the police. On the contrary, in blockchain systems, every transaction is effectively sent to the miners, which can be seen as a collective authority. This consequently lead to limited scalability because every transaction must be validated by the authority even when most of the transaction are legitimate.

Using the aforementioned insight, we explore an alternative consensus model for blockchain systems where transactions themself do not reach consensus, but nevertheless verifiable at a later stage by any node in the network. Informally, out model works as follows. Every node stores its own blockchain and every block is one transaction, same as the TrustChain construction. We randomly selected nodes in every round, the selected ones are called facilitators. They reach consensus not on the individual transactions, but on the state of every chain represented by a single digest, we call this state the checkpoint. If a checkpoint of some node is in consensus, then that node can proof to any other node that it holds a set of transactions that computes (form a chain) to the checkpoint. This immediately show that those transactions are tamper-proof.

We begin the detailed discussion by formulating the model Chapter 2. Next, we analyse the correctness, security and performance of our design in Chapter 3, this is where we present our theorems. Implementation and experimental results are discussed in Chapter 4. Finally, we compare our system with other state-of-the-art blockchain systems and conclude in Sections 5 and 6 respectively.

Chapter 2

System Architecture

The primary goal guiding our design is scalability. As mentioned in the Introduction, having a scalable blockchain system while still keeping global consensus allows the system to be ubiquitous and realise the full potential of blockchain.

The secondary goal is to design an application neutral system. In particular, it should act as a framework that provides the building blocks of blockchain based applications. Application developers using the framework should be able to create any application they wish. Further, we do not impose on a consensus algorithm, as long as it satisfies the properties of atomic broadcast which we describe in Section 2.1.2.

Due to the nature of our system, we do not explicitly address the Sybil attack [3]. Sybil defence mechanism always require some form of reputation score from the application. For example, social network based Sybil defence mechanims use graph structure of real-world relationships [9]. Online marketplaces such as Amazon use the rating of buyer and sellers. Thus it is not possible to design a Sybil defence mechanism with a a application neutral framework. On the other hand, our system also has no restrictions on the Sybil defence technique and application designers can pick the best mechanism for their application.

The third and final goal is security. Our system should be unaffected in the presence of powerful adversaries. Security is often difficult to verify, especially when it is not formalised, therefore we require our design to be provably secure. To summarise, our system design is designed with the following goals in mind.

- Application neutrality,

- scalability and

- security.

We begin the chapter with an intuitive overview of the architecture in Section 2.1. Next, we give the formal description, starting with the model

and assumptions in Section 2.2. Then, the three protocols which make up the complete system, namely consensus protocol (Section 2.3), transaction protocol (Section 2.4) and validation protocol (Section 2.5). Finally, the possible extensions are described in Section 2.6.

2.1 System Overview

The system consist of one data structure—Extended TrustChain, and three protocols—consensus protocol, transaction protocol and validation protocol. We first describe each component individually and then explain how they fit together in Section 2.1.4.

2.1.1 Extended TrustChain

Extended TrustChain naturally builds on top of TrustChain, thus we first describe the standard TrustChain. Our description has minor differences compared to the description in [**trustchain**]. This is to help with the description of the extended TrustChain. However, the two descriptions are functionally the same.

Standard TrustChain

In TrustChain, every node has a "personal" chain. Initially, the chain only contains a genesis block. When a node wishes to add a new transaction (TX), a new TX block is generated and is appended to the chain. A TX block must have a valid hash pointer pointing to the previous block and a reference[1] to its *pair*. As a result, a single transaction generates two TX blocks, one on each party's chain. An example of is shown in Figure 2.1.

If every node follows the rules of TrustChain and we only consider hash pointers, then the chain effectively forms a singly linked list. However, if a node violates the rules, then a *fork* may happen. That is, there may be more than one TX block with a hash pointer pointing back to the same block. In Figure 2.1, node b (in the middle chain) created two TX blocks that both point to $t_{b,5}$. If this is a ledger system it can be seen as a double spend, where the currency accumulated up until $t_{b,5}$ are spent twice.

Extended TrustChain

We are now ready to explain the Extended TrustChain, which we abbreviate to ETC. In ETC, we introduce a new type of block—checkpoint (CP) block. In contract to TX blocks, CP blocks do not store transactions or contain

---

[1]This is different from the original TrustChain definition found in [**trustchain**]. In there, a TX block has two outgoing edges which are hash pointers to the two parties involved in the transaction. This work uses one outgoing edge and a reference.

Figure 2.1: Every block is denoted by $t_{i,j}$, where i is the node ID and j is the sequence number of the block. Thus we have three nodes and three corresponding chains in this example. The arrows represent hash pointers and the dotted lines represent references. The blocks at the ends of one dotted line are pairs of each other. The red block after $t_{b,5}$ indicate a fork.

Figure 2.2: The circles represent CP blocks, they also have hash pointers (arrow) but do not have references (dotted line). Note that the sequence number counter do not change, it is shared with TX blocks.

references. Their purpose is to capture the state of the chain and the state of the whole system. In particular, the state of the chain is captured with a hash pointer. The state of the whole system is captured in the content of the CP block, namely as a digest of the latest *consensus result* which we explain in Section 2.1.2. A visual representation is shown in Figure 2.2.

2.1.2   Consensus Protocol

Before describing our consensus protocol, we take a brief detour to explain asynchronous common subset primitive and contrast it with other well known primitives. It is a fundamental building block of our consensus protocol.

Asynchronous common subset

TODO give history/background of ACS TODO more detailed explination with algorithm and new section. Asynchronous common subset (ACS) is an especially useful primitive for blockchain systems. It allows any party to

propose a value and the result is the set union of all the proposed values by the majority. Concretely, ACS needs to satisfy the following properties (adapted from [4]).

Definition 1. Properties of ACS

1. Validity: *If any correct party outputs a set C, then $|C| \geq n - t$ and C contains the input of at least $n - 2t$ parties.*

2. Agreement: *If a correct party outputs C, then every party outputs C.*

3. Totality: *If $n - f$ party receive an input, then all correct parties produce an output.*

ACS has the nice property of censorship resiliance when compared to other consensus algorithms. For instance, Hyperledger and Tender mint uses Practical Byzantine Fault Tolerance (PBFT) as their consensus algorithm. In PBFT, a leader is elected, if the leader is malicious but follows the protocol, then it can selectivly filter transactions. In contract, every party in ACS are involved in the proposal phase, and it guarantees that if $n - 2t$ parties propose the same transaction, then it must be in the agreed output.

The main drawback with ACS and other BFT protocols is the high message complexity. Typically, such protocols have the complexity of $O(n^2)$, where n is the number of parties. Hence, it may work with a small number of nodes, but it is infeasible for blockchain systems where thousands of nodes are involved.

Consensus Protocol

The consensus protocol runs continuously in rounds because a blockchain systems always need to reach consensus on new values, which are CP blocks in our case. This can be seen as running infinitely many rounds of some Byzantine consensus algorithm, starting a new execution immediately after the previous one is completed.

As we mentioned earlier, the high message complexity prohibits us from running a Byzantine consensus algorithm on a large network. Thus, for every round, we randomly select some node—called facilitators—to collect CP blocks and use them as the proposal. The facilitators are elected using a *luck value*, which is computed using $H(\mathcal{C}_r||pk_i)$, where $\mathcal{C}_r$ is the consensus result in round r and pk_i is the public key of i. Intuitively, the election is guaranteed to be random because the output of a cryptographically secure hash function is unpredictable and $\mathcal{C}_r$ cannot be determined in advance.

A visual explaination can be found in Appendix A, it walks through the steps needed for a node to be selected as a facilitator.

2.1.3 Transaction and Validation

The TX protocol is a simple request and response protocol. The nodes exchange one round of messages and create new TX blocks on their respective chains. Thus, as we mentioned before, one transaction should result in two TX blocks.

The consensus and transaction protocol by themselves do not provide a mechanism to detect forks or other forms of tamperaing. Thus we need a validation protocol to counteract malicious behaviour. When a node wish to validate one of its TX, it asks the counterparty for the *fragment* of the TX. A fragment of a TX is a section of the chain beginning and ending with CP blocks that contains the TX. Upon the counterparty's response, the node checks that the CP blocks are in consensus, the hash pointers are valid and his TX is actually in the fragment. The TX is valid if these conditions are satisfied. Intuitively, this works because it is hard (because hash collision is hard) to create a different chain that begins and ends with the same two CP blocks but with a different middle section.

2.1.4 Combined Protocol

The final protocol is essentially the concurrent composition of the three aforementioned protocols, all making use the Extended TrustChain data structure.

Our subprotocol design gives us the highly desireable non-blocking property. In particular, we do not need to "freeze" the state of the chain for some communication to complete in order to create a block. For instance, a node may start the consensus protocol, and while it is running, the node may still perform transactions. By the time the consensus protocol is done, the new CP block is added to whatever the state that the chain is in. It is not necessary to keep the chain immutable while the consensus protocol is running.

2.2 Model and Assumptions

For notational clarify, we use the following convention (adapted from [4]) throughout this work.

- Lower case (e.g. x) denotes a scalar object or a tuple.

- Upper case (e.g. X) denotes a set or a constant.

- Sans serif (e.g. $\mathsf{fn}(\cdot)$) denotes a function.

- Typewrite (e.g. `ack`) denotes message type.

Further, we use $a||b$ to dedote concatination of the binary representations of a and b.

We assume a distributed network where nodes are fully connected, the channels are reliable[2], but messages may be re-ordered and delayed by at most some time Δ, this is sometimes known as a Δ-synchronous network. We assume there exist a Public Key Infrastructure (PKI), and nodes are identified by their unique and permanent public key. Finally, we use the random oracle model, i.e. calls to the random oracle are denoted by $\mathsf{H} : \{0,1\}^* \rightarrow \{0,1\}^\lambda$, where $\{0,1\}^*$ denotes the space of finite binary strings and λ is the security parameter [1].

In our model we consider N nodes, which is the population size. n of them are facilitators, t out of n are malicious and the inequality $n \geq 3t+1$ must hold.

We use a restricted version of the adaptive corruption model. The first restriction is that corrupted node can only change across rounds. That is, if a round has started, the corrupted nodes cannot changed until the next round. The second restriction is that the adversary, presumably controlling all the corrupted nodes, is forgetful. Namely the adversary may learn the internal state such as the private key of a corrupted node, but if the node recovers, then the adversary must forget the private key. Otherwise the adversary can eventually learn all the private keys and sabotage the system. Finally, we assume computational security. That is, the adversary can run polynomial-time algorithms but not exponential-time algorithms.

The primary data structure used in our system is Extended TrustChain. Each node u has a public and private key pair—pk_u and sk_u, and a chain B_u. The chain consist of blocks $B_u = \{b_{u,i} : k \in \{0, \ldots, h-1\}\}$, where $b_{u,i}$ is the ith block of u, and $h = |B_u|$. We often use $b_{u,h}$ to denote the latest block. There are two types of blocks, TX blocks and CP blocks. If T_u is the set of TX blocks of u and C_u is the set of CP blocks of u, then it must be the case that $T_u \cup C_u = B_u$ and $T_u \cap C_u = \varnothing$. The notation $b_{u,i}$ is generic over the block type. We assume there exist a function $\mathsf{typeof} : B_u \rightarrow \{\tau, \gamma\}$ that returns the type of the block, where τ represents the TX type and γ represents the CP type.

2.2.1 TX Block

The TX block is a six-tuple, i.e $t_{u,i} = \langle \mathsf{H}(b_{u,i-1}), seq_u, txid, pk_v, m, sig_u \rangle$. We describe each item in turn.

1. $\mathsf{H}(b_{u,i-1})$ is the hash pointer to the previous block.

2. seq_u is the sequence which should equal i.

---

[2]Reliability can be achieved in unreliable networks by resending messages or using some error correction code.

3. *txid* is a cryptographically secure random number representing the transaction identifier.

4. pk_v is the public key of the counterparty.

5. m is the transaction message.

6. sig_u is the signature created using sk_u on the concatination of the binary representation of the five items above.

The fact that we have no constraint on the content of m is in alignment with our design goal—application neutrality.

TX blocks come in pairs. In particular, for every block $t_{u,i} = \langle \mathsf{H}(b_{u,i-1}), seq_u, txid, pk_v, m, sig_u \rangle$ there exist one and only one *pair* $t_{v,j} = \langle \mathsf{H}(b_{v,j-1}), seq_v, txid, pk_u, m, sig_v \rangle$. Note that the *txid* and m are the same, and the public keys refer to each other. Thus, given a TX block, these properties allow us to identify its pair.

2.2.2   CP Block

The CP block is a five-tuple, i.e. $c_{u,i} = \langle \mathsf{H}(b_{u,i-1}), seq_u, \mathsf{H}(\mathcal{C}_r), r, sig_u \rangle$, where $\mathcal{C}_r$ is the consensus result in round r, the other items are the same as the TX block definition. Note that unlike in our prior work [6], CP blocks and TX blocks do not have independent sequence numbers.

The genesis block in the chain must be a CP block in the form of $c_{u,0} = \langle \mathsf{H}(\bot), 0, \mathsf{H}(\bot), 0, sig_u \rangle$ where $\mathsf{H}(\bot)$ can be interpreted as applying the hash function on an empty string. The genesis block is unique due to every node due to sig_u.

2.2.3   Consensus Result

Our consensus protocol runs in rounds as discussed in Section 2.1. Every round is identified by a round number r, which is incremented on every new round. The consensus result is a tuple, i.e. $\mathcal{C}_r = \langle r, C \rangle$, where C is a set of CP blocks agreed by the facilitators of round r.

2.2.4   Chain Properties

Here we define a few important properties which results from the interleaving nature of CP and TX blocks.

If there exist a tuple $\langle c_{u,a}, c_{u,b} \rangle$ for a TX block $t_{u,i}$, where

$$a = \underset{k, k < i, \mathtt{typeof}(b_{u,k}) = \gamma}{\arg\min} (i - k)$$

$$b = \underset{k, k > i, \mathtt{typeof}(b_{u,k}) = \gamma}{\arg\min} (k - i),$$

then $\langle c_{u,a}, c_{u,b} \rangle$ is the *enclosure* of $t_{u,i}$. Some TX blocks may not have any subsequent CP blocks, then its enclosure is $\perp$.

If the enclosure of some TX block is $\langle c_{u,a}, c_{u,b} \rangle$, then its *fragment* is computed as $\{b_{u,i} : a \leq i \leq b\}$. For convenience, the function $\mathsf{fragment}(\cdot)$ represents the fragment of some TX block if it exists, otherwise $\perp$.

Agreed enclosure is the same as enclosure with an extra constraint where the CP blocks must be in some consensus result $\mathcal{C}_r$ and also must be the smallest enclosure. That is, suppose a chain is in the form $\{c_i, c_{i+1}, t_{i+2}, c_{i+3}\}^3$ and c_i, c_{i+1}, c_{i+3} are in $\mathcal{C}_r, \mathcal{C}_{r+1}, \mathcal{C}_{r+3}$ respectively, then the agreed enclosure of t_{i+2} is $\langle c_{i+1}, c_{i+3} \rangle$ and cannot be $\langle c_i, c_{i+3} \rangle$. Similarly, *agreed fragment* is computed using the agreed enclosure. We define its function to be $\mathsf{agreed\_fragment}(\cdot)$

The length of the fragment is constrained by L, namely $\forall t |\mathsf{fragment}(t)| \leq L$. The purpose to prevent spam and encourage nodes to create more CP blocks. L should be sufficiently high so that busy nodes are not hindered by it.

2.3   Consensus Protocol

Our consensus protocol runs on top of the model described above. It is directly related to the creation of CP blocks. The objectives of the protocol are to allow honest nodes always make progress (by creating new CP blocks), compute correct consensus result in every round and have unbiased election of facilitators. Concretely, we define the necessary properties as follows, it is collectively known as the correctness properties of the consensus protocol.

Definition 2. Correctness properties of the consensus protocol
   $\forall r \in \mathbb{N}$, *the following properties must be satisfied.*

   1. Agreement: *If one honest node outputs a list of facilitators $\mathcal{F}_r$, every other node honest decides on $\mathcal{F}_r$*

   2. Validity: *If any honest node outputs $\mathcal{F}_r$, then $|\mathcal{F}_r| > n$ and $|\mathcal{F}_r|$ must contain at least $n - t$ honest nodes.*

   3. Fairness: *Every node should have an equal probability of becoming a facilitator.*

   4. Termination: *At the end of the round, every node outputs some $\mathcal{F}_r$.*

Note that the final three properties are on the facilitators rather than on every node, they are the properties ACS. The fair lottery and consistent facilitator are properties on every node, they are prerequites for ACS.

We proceed to describe the protocol step by step. Starting with the bootstrap phase and then moving on to the actual consensus phase.

---

[3]Usually the notation is of the form $c_{u,i}$, but the node identity is not important here so we simplify it to c_i

2.3.1 Bootstrap Phase

Recall that facilitators are computed from the consensus result, but the consensus result is agreed by the facilitators. Thus we have a dependency cycle. The goal of the bootstrap phase is to give us a starting point in the cycle.

Imagine that there is some bootstrap oracle, that initiates the code on every node. The code satisfied all the properties in Definition 2. Namely every node has the same set of valid facilitators $\mathcal{F}_1$ that are randomly choosen. This concludes the bootstrap phase.

In practice, the bootstrap oracle is most likely the software developer and some of the desired properties cannot be achieved. In particular, it is not possible to have the fairness property because it is unlikely that the developer knows the identity of every node in advance.

2.3.2 Consensus Phase

The consensus phase begins when $\mathcal{F}_r$ is available to all the nodes. Note that $\mathcal{F}_r$ indicates the facilitators that were elected using result of round r and are responsible for driving the ACS protocol in round $r + 1$. The goal is to reach agreement on a set of new facilitators $\mathcal{F}_{r+1}$ that satisfied Definition 2.

There are two scenarios in the consensus phase. First, if node u is not the facilitator, it sends $\langle \texttt{cp\_msg}, c_{u,h} \rangle$ to all the facilitators. Second if the node is a facilitator, it waits for some duration D and collect messages of type $\texttt{cp\_msg}$, where $D \gg \Delta$. Invalid messages are removed, namely blocks with invalid signatures and duplicate blocks signed by the same key. After D elapses, it begins the ACS protocol and some $\mathcal{C}_{r+1}$ should be agreed upon by the end of it.

At the core of the consensus phase is the ACS protocol. While any ACS protocol that satisfies the standard definition will work, we use a simplification of HoneyBadgerBFT as our ACS protocol because it is the only (to the best of our knowledge) consensus algorithms designed for blockchain systems. We do not use the full HoneyBadgerBFT due to the following. First, the transactions in HoneyBadgerBFT are first queued in a buffer and the main consensus algorithm starts only when the buffer reaches an optimal size. We do not have an infinite stream of CP blocks, thus buffering is unsuitable. Second, HoneyBadgerBFT uses threshold encryption to hide the content of the transactions. But we do not reach consensus on transactions, only CP blocks, so hiding CP blocks is meaningless for us as it contains no transactional information.

When $\mathcal{F}_r$ reaches agreement on $\mathcal{C}_{r+1}$, they immediately broadcast two messages to all the nodes— first the consensus message $\langle \texttt{cons\_msg}, \mathcal{C}_{r+1} \rangle$, and second the signature message $\langle \texttt{cons\_sig}, r, sig \rangle$. The reason for sending $\texttt{cons\_sig}$ is the following. Recall that channels are not authenticated,

and there are no signatures in $\mathcal{C}_{r-1}$. If a non-facilitator sees some $\mathcal{C}_{r-1}$, it cannot immediately trust it because it may have been forged. Thus, To guarantee authenticity, every facilitator sends an additional message that is the signature of $\mathcal{C}_{r+1}$.

Continuing, upon receiving $\mathcal{C}_{r+1}$ and at least $n - f$ valid signatures by some u, u performs two asks. First, it creates a new CP block using new_cp$(\mathcal{C}_{r+1}, r + 1)$ (Algorithm 2). Second, it computes the new facilitators using get_facilitator(C, n) (Algorithm 1) and updates its facilitator list to the result. This concludes the consensus phase and brings us back to the beginning of the consensus phase.

Algorithm 1 Function get_facilitator(C, n) takes a list of CP blocks C and an integer n, sort evey element in C by its luck value (the λ-expression), and outputs the smallest n elements.

take$(n, \mathsf{sort\_by}(\lambda x.\mathsf{H}(x||pk \text{ of } x), C))$

Algorithm 2 Function new_cp$(\mathcal{C}_r, r)$ runs in the context of the caller u. It creates a new CP block and appends it to u's chain.

$h \leftarrow |B_u|$
$c_{u,h} \leftarrow \langle \mathsf{H}(b_{u,h-1}), h, \mathsf{H}(\mathcal{C}_r), r, sig_u \rangle$
$B_u \leftarrow B_u \cup c_{u,h}$

2.4 Transaction Protocol

The TX protocol, shown in Algorithm 4, is run by all nodes. It is also known as True Halves, first described by Veldhuisen [7, Chapter 3.2]. Node that wish to initiate a transaction calls new_tx$(pk_v, m, txid)$ (Algorithm 3) with the intended counterparty v identified by pk_v and message m. $txid$ should be a uniformly distributed random value, i.e. $txid \in_R \{0,1\}^{256}$. Then the initiator sends $\langle \mathtt{tx\_req}, t_{u,h} \rangle$ to v.

Algorithm 3 Function new_tx$(pk_v, m, txid)$ generates a new TX block and appends it to the caller u's chain. It is executed in the private context of u, i.e. it has access to the sk_u and B_u. The necessary arguments are the public key of the counterparty pk_v, the transaction message m and the transaction identifier $txid$.

$h \leftarrow |B_u|$
$t_{u,h} \leftarrow \langle \mathsf{H}(b_{u,h-1}), h, txid, pk_v, m, sig_u \rangle$
$B_u \leftarrow B_u \cup \{t_{u,h}\}$

Algorithm 4 The TX protocol which runs in the context of node u.

Upon $\langle \texttt{tx\_req}, t_{v,j} \rangle$ from v

    $txid, pk_v, m \leftarrow t_{v,j}$ $\triangleright$ unpack the TX

    $\texttt{new\_tx}(pk_u, m, txid)$

    store $t_{v,j}$ as the pair of $t_{u,h}$

    send $\langle \texttt{tx\_resp}, t_{u,h} \rangle$ to v

Upon $\langle \texttt{tx\_resp}, t_{v,j} \rangle$ from v

    $txid, pk_v, m \leftarrow t_{v,j}$ $\triangleright$ unpack the TX

    store $t_{v,j}$ as the pair of the TX with identifier $txid$

A key feature of the TX protocol is that it is non-blocking. At no time in Algorithm 3 or Algorithm 4 do we need to hold the chain state and wait for some message to be delivered before committing a new block to the chain. This allows for high concurrency where we can call $\texttt{new\_tx}(\cdot)$ multiple times without waiting for the corresponding $\texttt{tx\_resp}$ messages.

2.5 Validation Protocol

Up to this point, we do not provide a mechanism to detect forks or other forms of tampering or forging. The validation protocol aims to solve this issue. The protocol is also a request-response protocol, just like the transaction protocol. But before explaining the protocol itself, we first define what it means for a transaction to be valid.

2.5.1 Validity Definition

A transaction can be in one of three states in terms of validity—*valid, invalid* and *unknown*. Given a fragment $F_{v,j}$, the validity of the transaction $t_{u,i}$ is captured by the function $\texttt{get\_validity}(t, F)$ (Algorithm 5). The first four conditions (up to Line 21) essentially check whether the fragment is the one that the verifier needs. If it is not, then the verifier cannot make any decision and return *unknown*. This is likely to be the case for new transactions because $\texttt{agreed\_fragment}(\cdot)$ would be $\bot$. The next two conditions checks for tampering or missing blocks, if any of these misconducts are detected, then the TX is invalid.

Note that the validity is on a transaction (two TX blocks with the same $txid$), rather than on one TX block owned by a single party. It is defined this way because the malicious sender may create new TX blocks in their own chain but never send $\texttt{tx\_req}$ messages. In that case, it may seem that the counterparty, who is honest, purposefully ommitted TX blocks. But in reality, it was the malicious sender who did not follow the protocol. Thus, in such cases the whole transaction, identified by its $txid$ is marked as invalid.

Further, the caller of get_validity($t_uu, i, F_{v,i}$) is not necessarily u[4] Any node w may call get_validity($t_uu, i, F_{v,i}$) as long as the caller w has an agreed fragment of $t_{u,i}$—$F_{u,i}$. $F_{u,i}$ may be readily available if $w = u$ or it may be from some other vd_resp message, which we describe next in the validation protocol.

2.5.2 Validation Protocol

With the validity definition, we are ready to construct a protocol for determining the validity of transactions. The protocol is a simple response and request protocol (Algorithm 6). If u wishes to validate some TX with ID $txid$ and counterparty v, it sends $\langle$vd_req, $txid\rangle$ to v. The desired properties of the validation protocol are as follows.

Definition 3. Consensus on transactions

 1. Correctness: *The validation protocol outputs the correct result according to the aforementioned validity definition.*

 2. Agreement: *If any correct node decides on the validity (except when it is* unknown*) of a transaction, then all other correct nodes are able to reach the same conclusion or* unknown.

 3. Liveness: *Any valid transactions can be validated eventually.*

We make two remarks. First, just like the TX protocol, we do not block at any part of the protocol. Second, suppose some $F_{v,j}$ validates $t_{u,i}$, then that does not imply that $t_{u,i}$ only has one pair $t_{v,j}$. Our validity requirement only requires that there is only one $t_{v,j}$ in the correct consensus round. The counterparty may create any number of fake pairs in a later consensus rounds. But these fake pairs only pollutes the chain of v and can never be validated because the round is incorrect.

2.6 Protocol Extensions

(Move to future work?)

Up to this point, we have discussed our protocol in the context of the model and assumptions defined in Section 2.2. In this section, we remove a few assumptions and discuss how our architecture is adapted.

Gossip?

Churn?

Sybil?

---

[4]In practice it often is because after completing the TX protocol the parties are incentivised to check that the counterparty "did the right thing".

Algorithm 5 Function get_validity($t_{u,i}, F_{v,j}$) validates the transaction $t_{u,i}$. $F_{v,j}$ is the corresponding fragment received from v.

We assume there exist a valid $F_{u,i}$, namely the agreed fragment of $t_{u,i}$. The caller is w, it may be u but this is not necessary.

```
 1: c_{v,a} ← first(F_{v,j})
 2: c_{v,b} ← last(F_{v,j})
 3: if c_{v,a} or c_{v,b} are not in consensus then
 4:     return unknown
 5: end if                                    ▷ v has agreed fragment
 6:
 7: if |F_{v,j}| > L then
 8:     return unknown
 9: end if                                    ▷ fragment not too long
10:
11: if sequence number in F_{v,j} is correct (sequential) then
12:     if hash pointers in F_{v,j} is wrong then
13:         return unknown
14:     end if
15: end if                                    ▷ correct TrustChain structure
16:
17: c_{u,b} ← last(F_{u,i})
18: if c_{u,b} is not created using the same C_r as c_{v,b} then
19:     return unknown
20: end if                                    ▷ correct consensus round
21:
22: txid, pk_v, m ← t_{u,i}
23: if number of blocks of txid in F_{v,j} ≠ 1 then
24:     return invalid
25: end if                                    ▷ TX exists
26:
27: txid', pk'_u, m' ← t_{v,j}
28: if m ≠ m' ∨ pk_u ≠ pk'_u ∨ |F_{v,j}| > L then
29:     return invalid
30: end if                                    ▷ no tampering
31:
32: return valid
```

Algorithm 6 Validation protocol

Upon $\langle \mathtt{vd\_req}, txid \rangle$ from v
   $t_{u,i} \leftarrow$ the transaction identified by $txid$
   $F_{u,i} \leftarrow \mathsf{agreed\_fragment}(t_{u,i})$
   send $\langle \mathtt{vd\_resp}, txid, F_{u,i} \rangle$ to v
Upon $\langle \mathtt{vd\_resp}, txid, F_{v,j} \rangle$ from v
   $t_{u,i} \leftarrow$ the transaction identified by $txid$
   set the validity of $t_{u,i}$ to $\mathsf{get\_validity}(t_{u,i}, F_{v,j})$

Chapter 3

Fault Tolerance and Security Analysis

Up to this point we described our system specification in detail. Of course, specification along does not establish any truths. In this chapter, we prove two aspects of our system. First is correctness, where we show that the consensus protocol and the validation protocol satisfies their desired properties (Definition 2 and Definition 3 respectively). The second is performance, where we prove the lower bound of our throughput and show that it out performs classical blockchain systems.

3.1 Correctness in the Presense of Faults

Our first objective in this section is to establish truths regarding the correctness of our protocol. We do this in two parts. First we use mathematical induction to show that properties in Definition 2 holds for all round. Building on top that, if Definition 2 is true, then we can show that many properties in Definition 3 is true.

3.1.1 Correctness of Consensus

We begin our analysis by establishing the fact that the get_facilitator($\cdot$) is fair.

Lemma 1. *Every node with a CP block in $\mathcal{C}_r$, should have an equal probability to be elected as a facilitator.*

Proof. This directly follows from the random oracle model. Recall that the luck value is computed using $\mathsf{H}(\mathcal{C}_r, ||pk_i)$. Since pk_i is unique for every node that has a CP block in $\mathcal{C}_r$, the output of $\mathsf{H}(\cdot)$ is uniformly random. This implies that the ordered sequence by luck value is uniformly random. $\square$

Using Lemma 1, we show that our consensus protocol satisfied Definition 2.

Lemma 2. $\forall r \in \mathbb{N}$, *agreement, validity, fairness and termination holds.*

Proof. We proof by mathematical induction.

In the base case, agreement, validity fairness and termination follows directly from the bootstrap protocol, due to the bootstrap oracle. Note that the result is $\mathcal{F}_1$, which indicates the facilitators that are agreed in round 1 and are responsible for driving the ACS protocol in round 2.

For the inductive step, we assume that the properties hold for round r. Then, to start round $r + 1$, the honest nodes begin sending CP blocks to $\mathcal{F}_r$. Since the honest region in $\mathcal{F}_r$ waits for D and $D \gg \Delta$, the CP blocks of the honest nodes are guaranteed to be received by at least $n - t$ facilitators. The agreement property of ACS (from Definition 1) ensures that the $\mathcal{C}_{r+1}$ is in consensus. Observe that $\mathcal{F}_{r+1}$ is computed using the deterministic function get_facilitators($\cdot$). Thus agreement of $\mathcal{F}_{r+1}$ follows directly from the agreement of ACS. The validity property of ACS ensures that $\mathcal{C}_{r+1}$ contains the input of at least $n - 2t$ parties, but this is a quorum containing at least one honest facilitator, thus $\mathcal{C}_{r+1}$ contains the CP blocks of all nodes. Due to the assumption that the adversary cannot corrupt more than t nodes, validity of $\mathcal{F}_{r+1}$ also follows from validity of ACS. Since all honest nodes are in $\mathcal{C}_{r+1}$, the fairness property follows directly from Lemma 1. Finally, the termination property holds because D eventually elapses and then ACS eventually terminates (the totality property of ACS). This completes the proof. $\square$

3.1.2 Correctness of Validation

The consensus protocol (on CP blocks and facilitators) is the backbone for consensus on transactions. In this section we build on top of Lemma 2 to show that most properties in Definition 3 can be satisfied.

Lemma 3. *The validation protocol outputs the correct result according to the validity definition.*

Proof. The algorithm (Algorithm 5) is the validity definition. $\square$

Theorem 1. *If any correct node decides on the validity (except when it is* unknown*) of a transaction, then all other correct nodes are able to reach the same conclusion or* unknown*.*

Proof. We proof by contradiction. Without loss of generality, for some transaction t with an agreed fragment F, node u decides *valid* but node v decides *invalid*. Then there exist a fragment $F' = \{\ldots, t', c'\}$ which u received that contains a valid pair of t—t'. There also exist a fragment

$F'' = \{\ldots, t'', c''\}$ which v received that does not contain or contains an invalid pair—t''. In both cases, the get_validity$(\cdot)$ must have reached Line 21. Due to Lemma 2, we have $c' = c''$, otherwise the result would be *unknown*. Since $c'(= c'') = \langle H(t'), \ldots \rangle$ we must have $H(t') = H(t'')$ and $t' \neq t''$ (because t'' is invalid). In other words, the sender of F'' must be able to create some t'' that has the same digest as t'. But we assumed that the adversary can only perform polynomial-time algorithm, but in order to find t'' it needs to query the random oracle exponentially many times. Thus we have a contradiction and this completes our proof. $\square$

Theorem 1 is our first major result. It shows that consensus on CP blocks would lead to consensus on TX blocks when the nodes are running the validation protocol. One of the main advantages over running a consensus algorithm on all the transactions is that the rate of transaction is no longer dependent on the consensus algorithm—ACS (up to L transactions). This enables horizontal scalability where adding new nodes would lead to higher global transaction rate. In addition, a convenient consequence Theorem 1 is unforgeability. That is, no polynomial time adversary is able to create two chains $F = \{\ldots, t, c\}$ $F' = \{\ldots, t', c\}$ with correct hash pointers and the same end of chain c.

However, not everything is perfect. Now we show a negative result, where the liveness property cannot be attained.

Lemma 4. *There exist a valid transactions that cannot be validated eventually.*

Proof. We proof by providing a counterexample. Suppose nodes u and v correctly perfomed the TX protocol which resulted a transaction t. Then when u wants to validate t, it does so by sending vd_req message to v. v can act maliciously and ignore all vd_req message, thus t can never be validated. $\square$

Although this is a negative result, it does not put the adversary in an advantageous position. If the adversary is observed to ignore validation requests, then the honest nodes may prefer not to transact with her in the future. Thus, to stay relevant in the system, the adversary need to comply to the protocol.

3.2 Performance

TODO intro

3.2.1 Message Complexity of ACS

The message complexity of ACS is $O(n^2|v| + \lambda n^3 \log n)$ [4], where $|v|$ is the size of largest message and λ is the security parameter. Note that the se-

curity parameter is the same as the one for our random oracle described in Section 2.2. In particular, it is from the use of $H(\cdot)$ in the reliable broadcast phase in ACS. In our system, we wish to understand the scalability properties. Thus we consider the complexity as a function of N rather than n or λ. Since $|v|$ is at most all the CP blocks from every node, we have $|v| = N$. Therefore the message complexity of ACS in our system is $O(N)$. Since we use a constant n, $O(N)$ message complexity also holds for a single facilitator.

3.2.2 Bandwidth Requirement for Transactions

To make and validate a transaction, the bandwidth required is of $O(l)$, where l is the length of the agreed fragment. This can be seen from the fact that the largest message by far is the `vd_resp` message, which contains the agreed fragment, the other messages (`tx_req`, `tx_resp` and `vd_req`) are constant factors. If we assume that every node performs transactions at a constant rate of r_{tx} per second. Then $l = (D_{acs} + D) \cdot r_{tx}$, where D_{acs} is the duration an instance of ACS. But from Section 3.2.1, we know that D_{acs} is of $O(N)$, thus the bandwidth per transaction is $O(N)$. This is natural because consensus duration would be longer if there are more CP blocks, which means that the agreed fragments are longer. This behaviour is also verified experimentally in Chapter 4.

3.2.3 Global Throughput

Suppose every node has some fixed bandwidth capacity C (unit of communication per second), they make transactions at r_{tx} per second. Then we have the inequality $C \geq r_{tx}l$, where l is the length of the the agree fragment as before. Rearranging, we get $\frac{C}{l} \geq r_{tx}$. With this, we consider two cases, first is when all the nodes are running at maximum capacity. Recall that l is of $O(N)$, so as the the population increases, the transaction rate must decrease in order to maintain the inverse relationship, thus r_{tx} is of $O(N^{-1})$. Therefore, the global throughput would be $O(N^{-1})N = O(1)$. In the second case, nodes are not running at maximum capacity and r_{tx} is maintained. Therefore, if every node runs at r_{tx}, the global throughput becomes $O(N)$ until N is too large and we go back to the first case.

The upshot of this analysis is that our system scales nicely at a global throughput of $O(N)$ until the population gets too large. Then we maintain a constant global throughput. This result falls a bit short of what we envisioned in the introduction. However, the population is not the number of nodes that use the system, but the number of nodes that are online during a single round. Furthermore, this result is for the worst case where every transaction needs an agreed fragment to be transmitted. In practice, nodes are able to cache agreed fragments. For instance, if u and v make x transac-

22

tions in a single round, then only one agreed fragment need to be exchanged as it contains all the transactions rather than x agreed fragments.

3.3   Effect of A Highly Adverserial Environment

Chapter 4

Implementation and Experimental Results

Thus far, we have only discussed our system from a theoretical perspective. Henceforth, we evaluate out system experimentally and compare the results with the theoretical analysis. We begin this chapter by a description of the implementation in Section 4.1. Then, we move on to describing our experimental setup in Section 4.2. Finally, Section 4.3 presents our experimental results and our evaluation. Our experiment primarily focuses on the consensus duration and the throughput.

4.1   Implementation

The prototype implementation is done in the event driven paradigm, using the Python programming language. We use Twisted as our networking library. The code can be found on GitHub[1].

The structure of the implementation is primarily made up of three modules—`acs`, `trustchain` and `node`. `acs`, as its name suggests, implements ACS. We use liberasurecode[2] as the Reed-Solomon error correcting code library, used in RBC. An implementation detail is that liberasurecode cannot create more than 32 code blocks[3]. The `acs` module provides a small interface to the caller to start and stop the consensus process and also receive results. The `trustchain` module implements the Extended TrustChain data structure. It also provides the essential algorithm necessary to interact with Extended TrustChain such as new_tx($\cdot$), new_cp($\cdot$), agreed_fragment($\cdot$) and so on. Since this is a prototype implementation, we only store the data struc-

---

[1] `https://github.com/kc1212/consensus-thesis-code`

[2] `https://github.com/openstack/liberasurecode`

[3] The 32 code blocks limitation is hardcoded in the source file, see `https://github.com/openstack/liberasurecode/blob/0794b31c623e4cede76d66be730719d24debcca9/include/erasurecode/erasurecode.h#L35`

ture in memory and not on disk. Finally, the `node` module ties everything together. It implements the the consensus phase, the transaction protocol and the validation protocol.

Every node keeps a persistent TCP connection with every other node. This creates a fully connected network for our experiment. It is certainly not idea in real world scenarios where nodes may have limited resources. But as a prototype, it is sufficient to create a system that has just over a thousand nodes, which is enough for us to experiment on.

4.2   Experimental Setup

There are two asppects of the experimental setup. First is the nature of the experiment. Second is the physical setup. We describe these in turn.

The nature and the goal of the experiment is to run the three protocols—consensus protocol, transaction protocol and validation protocol—simultaneously and analyse the throughput and consensus duration. There are two types of parameters which we must consider. First are the fixed parameters r_{tx} and D. These are selected from emperical evidence, we found that $D = 30$ seconds to be more than enough for all CP blocks to reach consensus. We also found that using $r_{\text{tx}} = 8$ gave us good throughput without putting too much demand on the bandwidth which is must also be used for ACS. These parameters are fixed for all our experiments. The second set of parameters can be seen as the domain, and we run our experiment for every combination of these parameters. Concretely, there are three of them. The number of facilitators $n\ in\{8, 12, \ldots, 32\}$. The population size $N \in \{100, 200, \ldots, 1200\}$. Finally the two modes of transaction. The first mode or the ideal mode is that nodes only transact with their immediate neighbour. This minimises the bandwidth required per validated transaction because agreed fragment can be cached. The second mode is in the other extreme, where every transaction is with a random node out of the N nodes in the system, thus the agreed fragment is unlikely to be cached. Unfortunately, the maximum n is 32 because the limitation in liberasurecode mentioned in Section 4.1. N stops at 1200 is due to our physical setup, which we describe next.

The experiment is run on the DAS-5 (The Distributed ASCI Supercomputer 5). From now on, we use "machines" to refer to DAS-5 nodes and nodes to refer to a running instance in our system. On DAS-5 we use up to 24 machine, for each machine we use 50 nodes. This gives us the aforementioned 1200 number. With this setup, we cannot run more nodes because the every machine only has 65535 ports available (some of them are reserved). But 50 nodes each need 1200 TCP connections which is 60000 TCP connections per machine. Thus we set N to be at most 1200[4].

---

[4]While it is possible to have many more TCP connections per machine, but it requires additional network interface which is something we do not control on the DAS-5.

To coordinate nodes on many different machine, we employ a discovery server to inform every node the IP addresses and port numbers of every other node. It is only run before the experiment and is not used during the experiment.

4.3 Evaluation

4.3.1 Consensus Duration

4.3.2 Global Throughput

4.4 Evaluation

- How fast is the consensus algorithm? Possibly plot graph of time versus the number of nodes.

- Does the promoter registration phase add a lot of extra overhead?

- What's the rate of transaction such that they can be verified "on time", i.e. without a growing backlog?

- Our global validation rate is somewhat equivalent to the transaction rate in other systems. Does the validation rate scale with respect to the number of nodes? In theory it should. Plot validation rate vs number of nodes, we expect it to be almost linear.

Chapter 5

Related Work

Chapter 6

Conclusion

Bibliography

[1] Mihir Bellare and Phillip Rogaway. "Random oracles are practical: A paradigm for designing efficient protocols". In: *Proceedings of the 1st ACM conference on Computer and communications security*. ACM. 1993, pp. 62–73.

[2] Jason Bloomberg. *Eight Reasons To Be Skeptical About Blockchain*. 2017. URL: https://www.forbes.com/sites/jasonbloomberg/2017/05/31/eight-reasons-to-be-skeptical-about-blockchain/ (visited on 12/06/2017).

[3] John R Douceur. "The sybil attack". In: *International Workshop on Peer-to-Peer Systems*. Springer. 2002, pp. 251–260.

[4] Andrew Miller et al. "The honey badger of BFT protocols". In: *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security*. ACM. 2016, pp. 31–42.

[5] Justin OConnell. *BitTorrent Creator Cohen Talks Bitcoin, Blockchain*. 2017. URL: https://www.cryptocoinsnews.com/bram-cohen-talks-bitcoin/ (visited on 09/06/2017).

[6] Zhijie Ren et al. *Implicit Consensus: Blockchain with Unbounded Throughput*. 2017. eprint: arXiv:1705.11046.

[7] Pim Veldhuisen. "Leveraging blockchains to establish cooperation". MA thesis. Delft University of Technology, May 2017. URL: http://resolver.tudelft.nl/uuid:0bd2fbdf-bdde-4c6f-8a96-c42077bb2d49.

[8] *Visa Inc. at a Glance*. 2015. URL: https://usa.visa.com/dam/VCOM/download/corporate/media/visa-fact-sheet-Jun2015.pdf (visited on 12/06/2017).

[9] Haifeng Yu et al. "Sybilguard: defending against sybil attacks via social networks". In: *ACM SIGCOMM Computer Communication Review*. Vol. 36. 4. ACM. 2006, pp. 267–278.

Appendix A

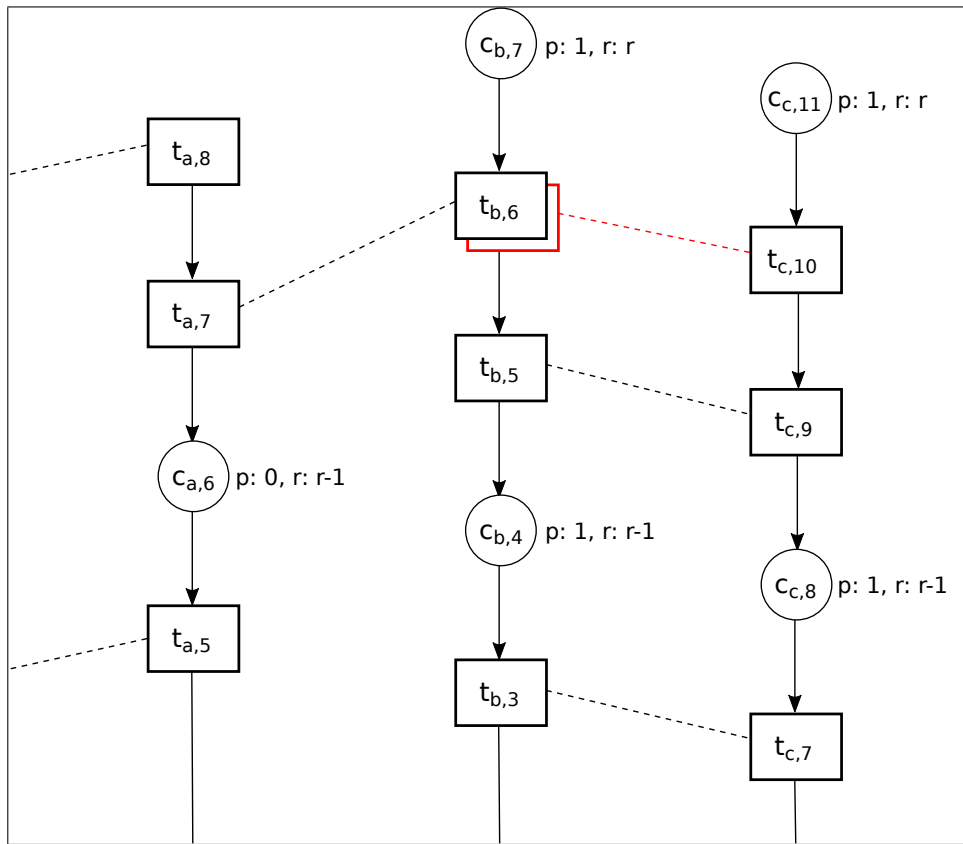# Consensus Example

Figure A.1: Initial state

Figure A.2: Initial state

Figure A.3: Initial state

Figure A.4: Initial state